Opened: Tuesday, 18 November 2025, 4:00 PM
Due: Monday, 1 December 2025, 11:59 PM

- Task 1: Obtain and prepare the data
- Task 2: Task 2: Build a plain autoencoder network, use it for embedding mages of digits
- Task 3: Perform cross-validation and model selection

Task 1: Obtain and prepare the data

Download the mnist data set. See Assignment 2, and recall that you can simply do:

```
from tensorflow.keras.datasets import mnist
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In the data preprocessing section of your code, select two classes among the available ten (0 1 2 3 4 5 6 7 8 9). Make this selection parametric so that you can perform experiments with different pairs, e.g.

```
class0 = 5; class1 = 7
```

or

```
classes = (4, 2)
```

(or whatever you like, how you store classes is irrelevant and the above are just suggestions, not prescriptions).

Task 2: Build a plain autoencoder network, use it for embedding mages of digits

Here you have to create a shallow neural network. It should have:

- As many output units as input units (for use with the MNIST data, that means d = 784)
- Two (2) hidden units
- Sigmoid activation on the hidden layer
- Linear activation on the output layer

Train the network with this configuration on the MNIST data, using X_train in place of y_train, and MSE as the loss.

Hints:

- You may have to use a small learning rate and do more than one learning trial to obtain convergence to a low local minimum.
- If you are using Keras, the easiest way to be able to use just a part of the network (take the output at the hidden layer, not using the output layer) is to use the .

Task 3: Perform cross-validation and model selection

Here you have to create another program, based on the previous one. It should:

- Select a list of possible values of *h*, and for each:
  - Run k-fold cross-validation and obtain *k* estimates of quality
  - For each "fold", repeat training several times (multi-start approach) and collect the best result, saving it in a list
  - Perform statistics on the list and store the result in *another* list, as a typical `mse` value and a typical range `[mselower, mseupper]` (e.g. median, 25th percentile, 75th percentile)
- Select the best value for *h* based on the cross-validated test results.

Remark: Note that this is a multi- (two-)criterion optimisation: we want the best typical mse (low error), but also the narrowest range (good generalisation).

This is an instance of the bias/variance dilemma. As such, there is no hard-and-fast rule to decide. Explain your choice.

Hints:

- Note that your code will have a certain number of nested loops: outermost, on values of h; then on k "folds"; then on a number of restarts.
- Larger errors may have larger variability. It may then be advisable to compare not directly the range extension, but the ratio `(mseupper-mselower)/msetypical`. This is equivalent to reasoning in log scale. Take care of possible numerical issues (e.g. what if typical mse = 0?)
- For each value of h, store all the histories to plot. Decide afterward, by inspection, which ones are interesting and should be plotted.

Add submission

Submission status

Attempt number	This is attempt 1.